

Generative adversarial networks

Lecture 9a

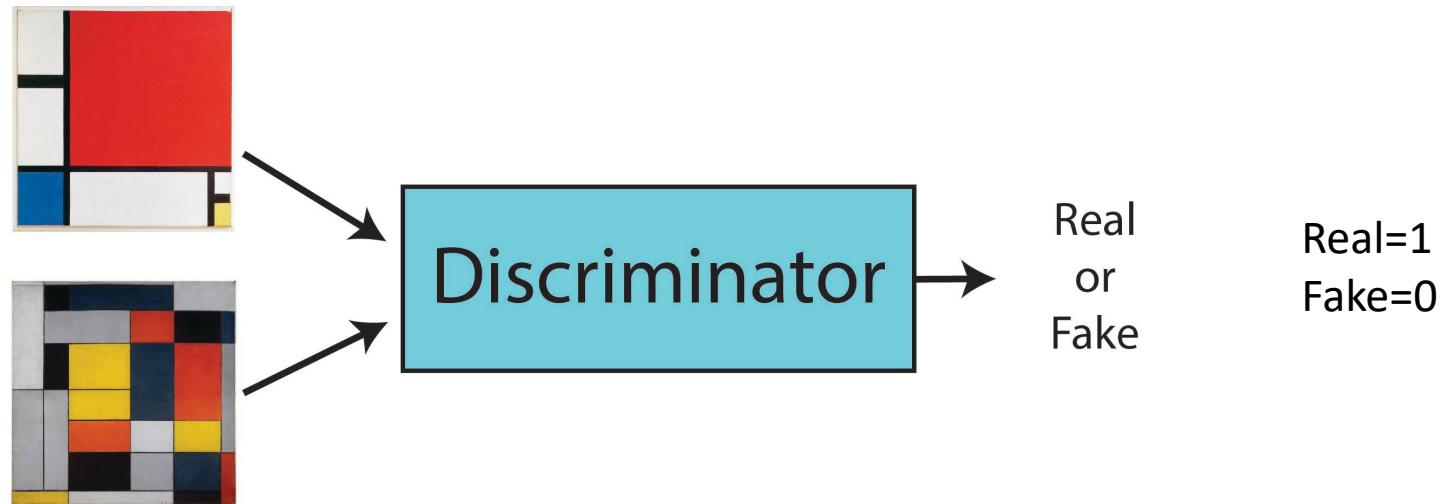
Fake images



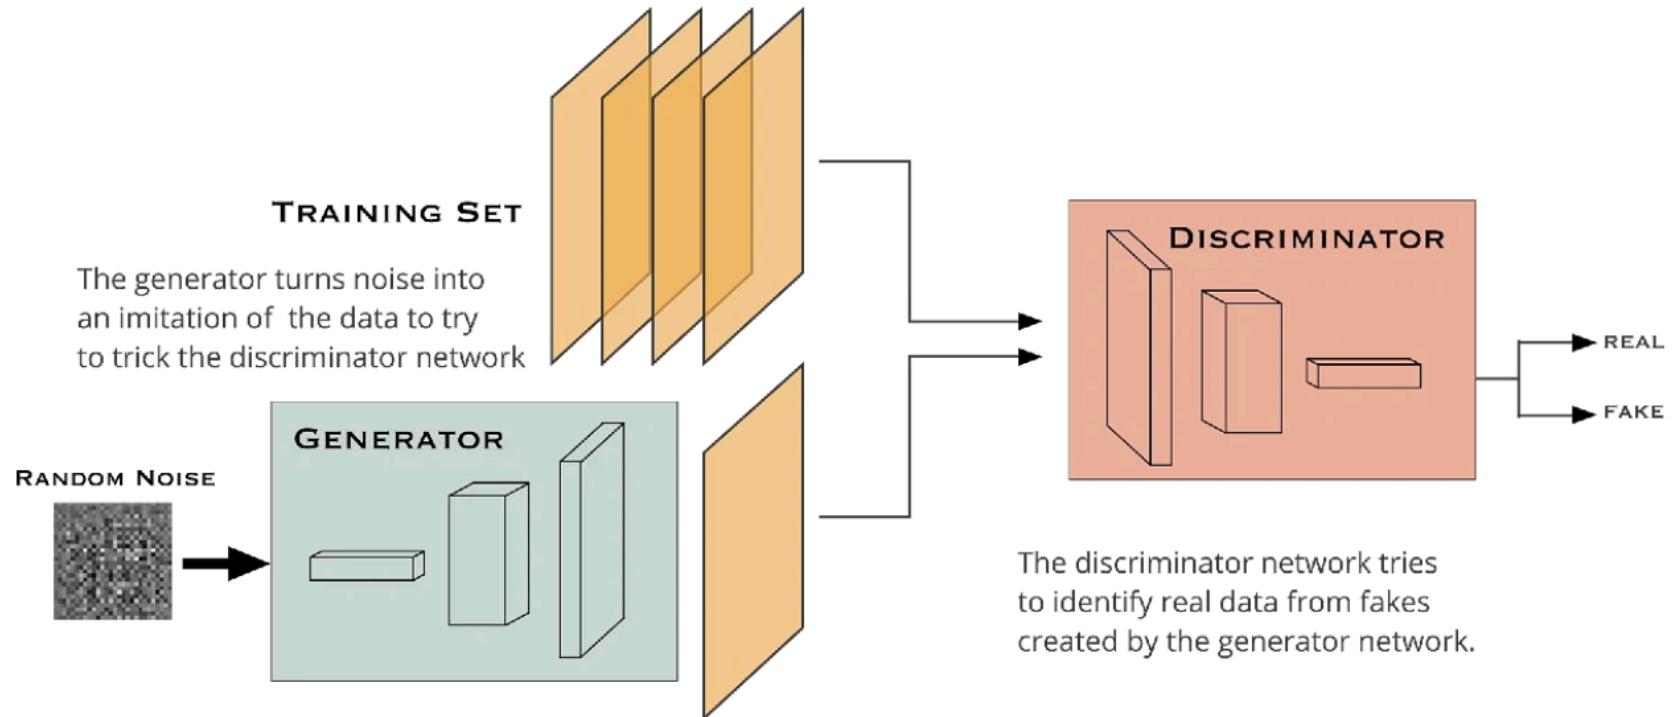
Discriminator



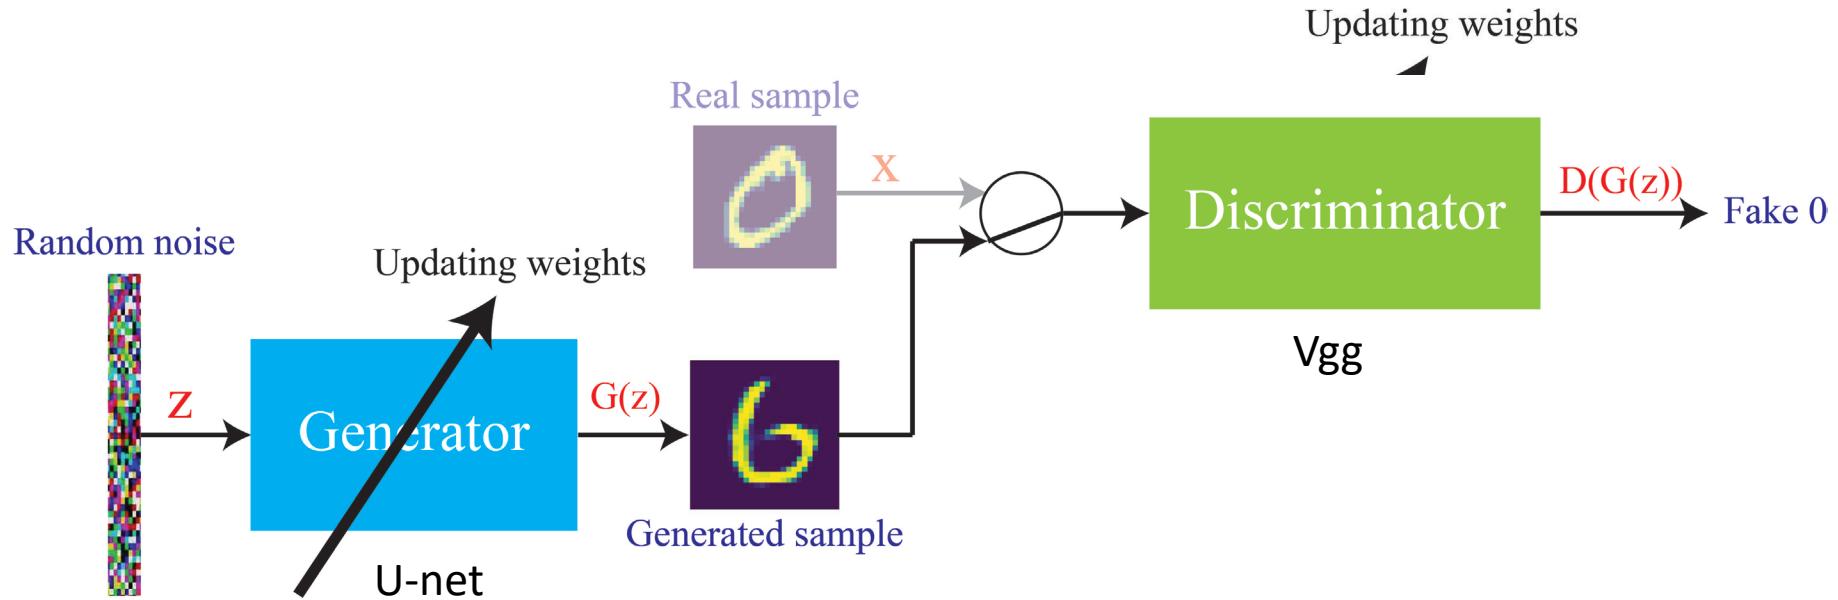
Mondrian paintings



Generative Adversarial Network (GAN)



Generative Adversarial Networks (GAN)



Cost function:

$$C = \sum_x [\log(D(x))] + \sum_z [\log(1 - D(G(z)))]$$

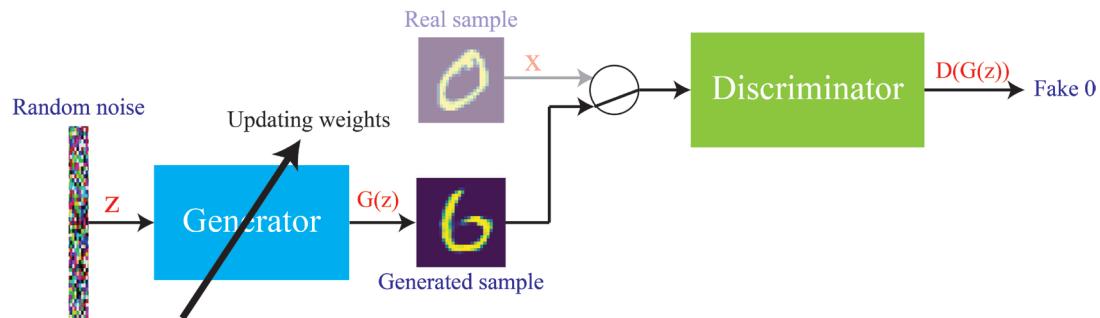
The **generator** tries to **minimize** the cost function

The **discriminator** tries to **maximize** the cost function

- x is the real data
- z is the random noise
- $D(x)$ is the discriminator output
- $G(z)$ is the generated data

Extreme behaviour (D= 0 or 1)

Real input	D=1	C=logD=0	DN no change	GN no change
Real input	D=0	C=log(0)=-∞	DN changes	GN no change
Fake input	D=1	C=log(1-D)=−∞	DN changes	GN no change
Fake input	D=0	C=log(1-D)=0	DN no change	GN changes



$$C = \sum_x \left[\log(D(x)) \right] + \sum_z \left[\log(1 - D(G(z))) \right]$$

GAN: generating digits

Generator

Layer	Size out	Activation	Kernel
Input	100		
Fully connected	100→12544	LeakyReLU	
Reshape	7x7x256		
Convolution	7x7x128	LeakyReLU	5x5
Up-sampling	14x14x128		
Convolution	14x14x64	LeakyReLU	5x5
Up-sampling	28x28x64		
Convolution	28x28x1	tanh	5x5

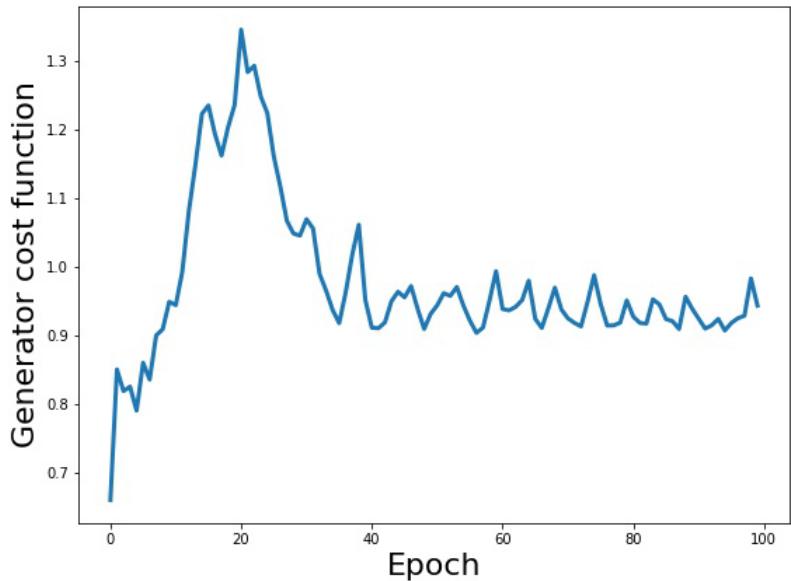
Discriminator

Layer	Size out	Activation	Kernel
Input	28x28x1		
Convolution	28x28x64	LeakyReLU	5x5
Down-sampling	14x14x64		
Dropout			
Convolution	14x14x128	LeakyReLU	5x5
Down-sampling	7x7x128		
Dropout			
Fully connected	6272→1		

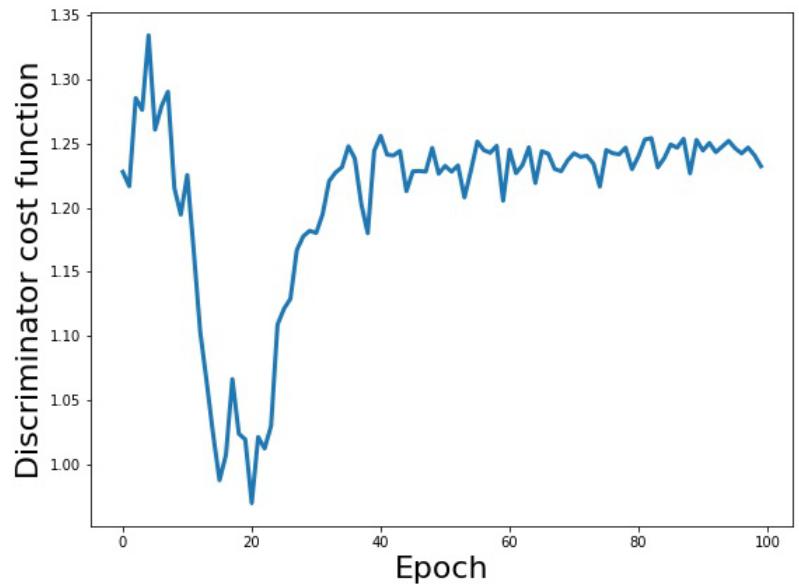
- Mnist digits as real samples
 - 60000 real samples
- Generator
 - Adam optimizer
 - Learning rate=0.0001
- Generator input:
 - Random normal between 0 and 1
- Discriminator
 - Adam optimizer
 - Learning rate=0.0001
- 100 epochs

GAN: generating digits

Generator

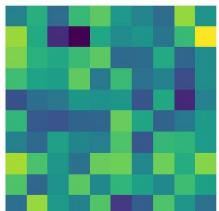


Discriminator

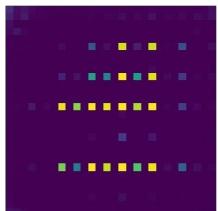


Generator

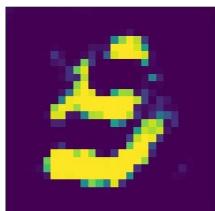
input



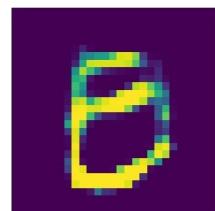
Epoch 1



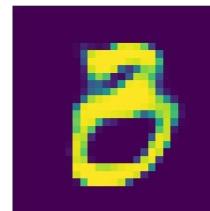
Epoch 10



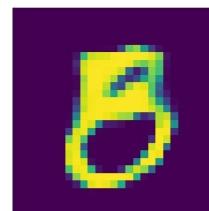
Epoch 30



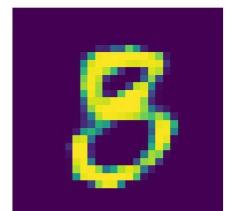
Epoch 60



Epoch 80



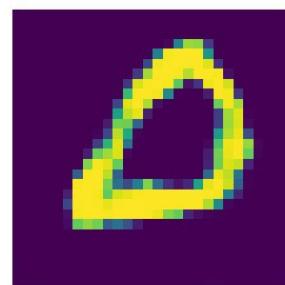
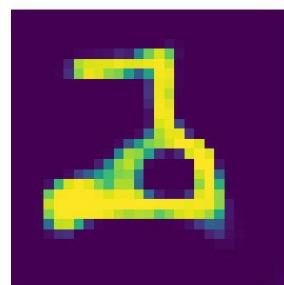
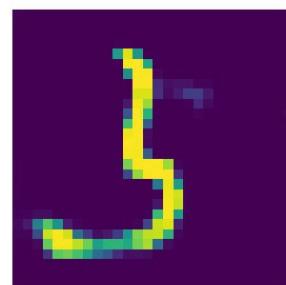
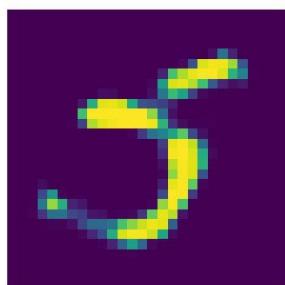
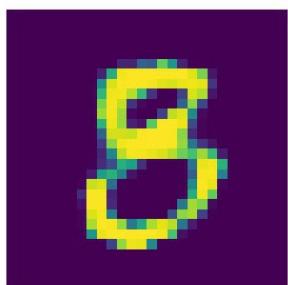
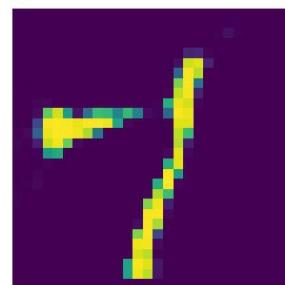
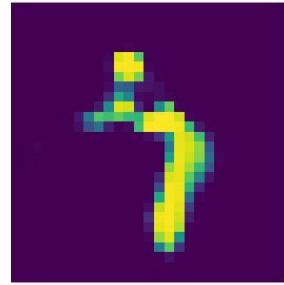
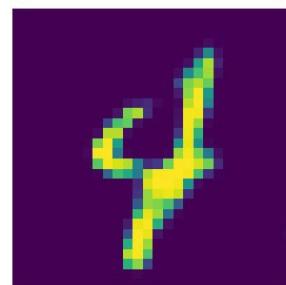
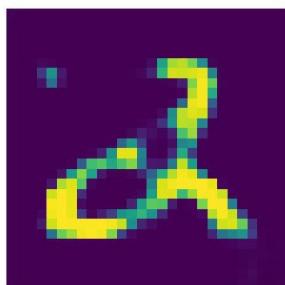
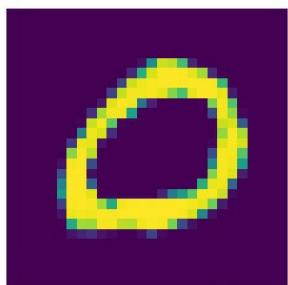
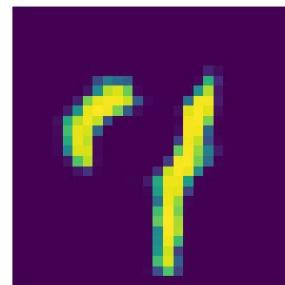
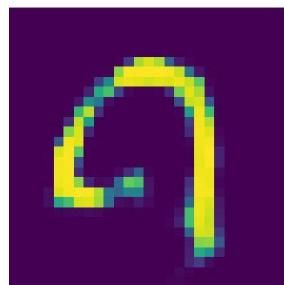
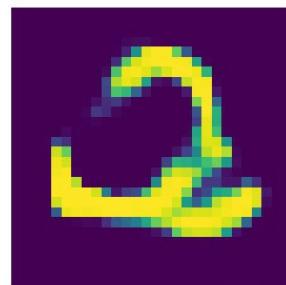
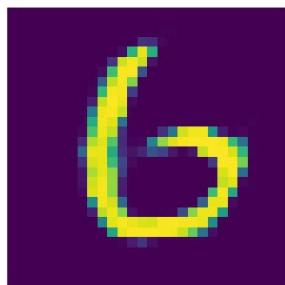
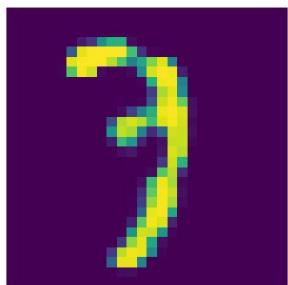
Epoch 100



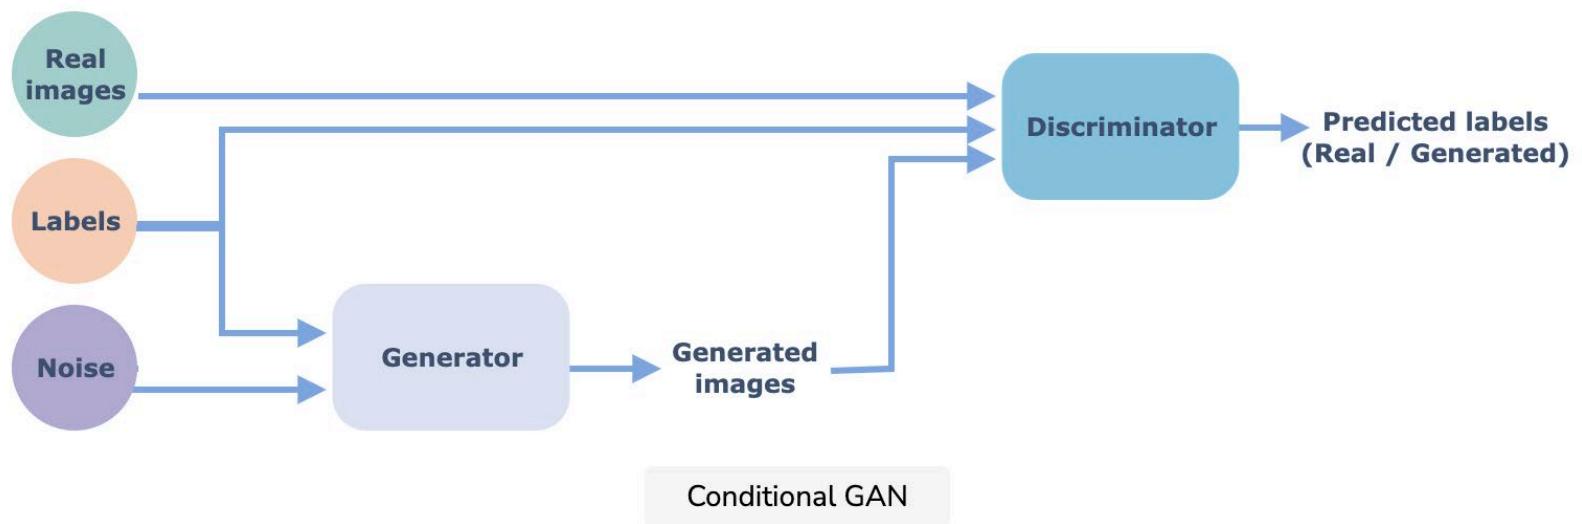
GAN: generating digits

Examples of generated digits

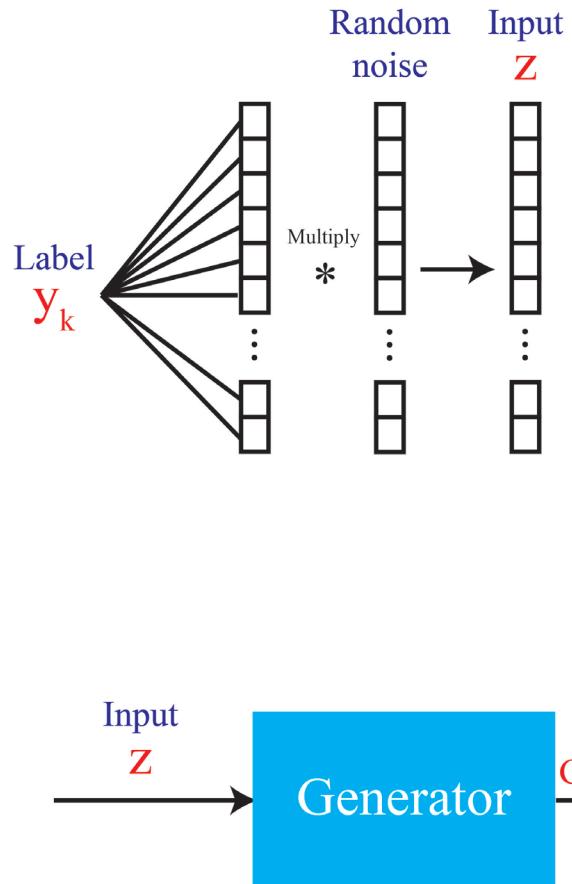
There is no control over which digit is going to be generated



Conditional GAN



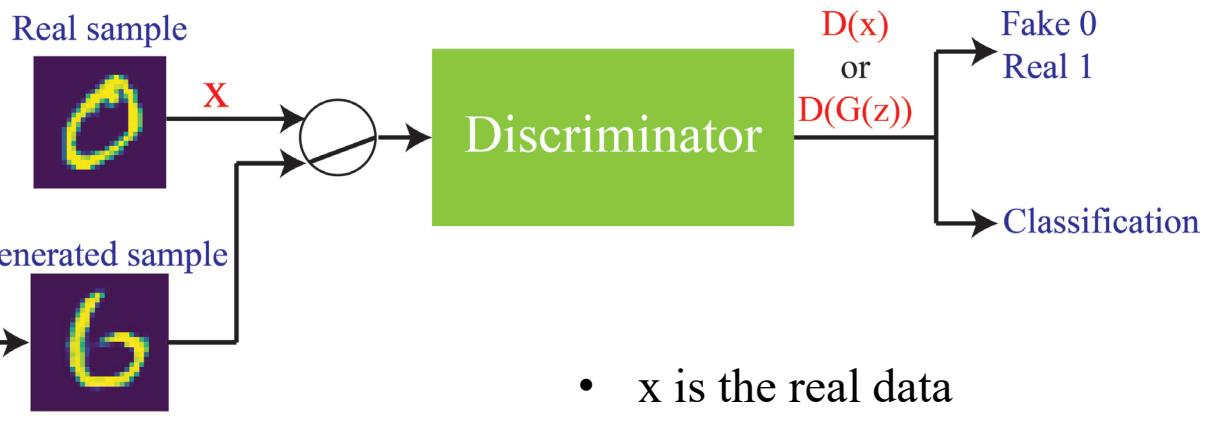
Auxiliary classifier GAN: AC-GAN



Cost function:

$$C_{\text{discrimination}} = \sum_x [\log(D(x))] + \sum_z [\log(1 - D(G(z)))]$$

$$C_{\text{classification}} = \sum_x \left[\sum_k (y_k^{(x)} \log(D(x))) \right] + \sum_z \left[\sum_k (y_k^{(z)} \log(D(G(z)))) \right]$$



- Discriminator maximizes $C_{\text{classification}} + C_{\text{discrimination}}$
- Generator maximizes $C_{\text{classification}} - C_{\text{discrimination}}$

- x is the real data
- z is the random noise
- $G(z)$ is the generated data
- $D(x)$ is the discriminator output
- $y_k^{(x)}$ is the label of sample x

AC-GAN: generating digits

Generator

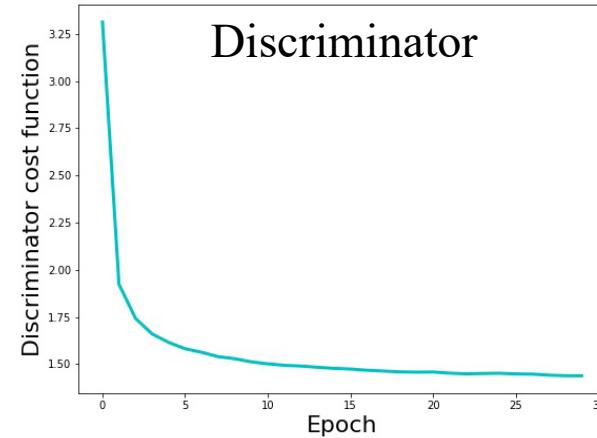
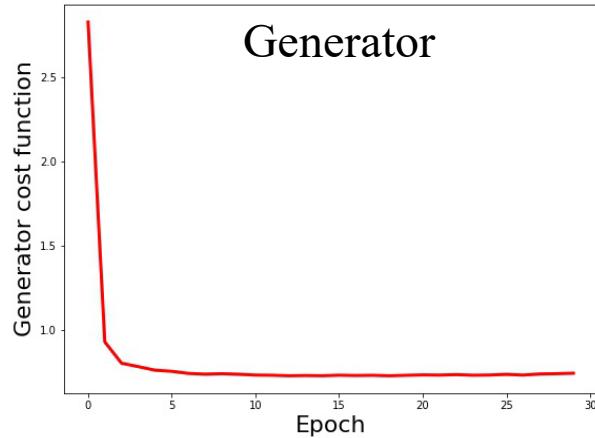
Layer	Size out	Activation	Kernel
Input	100		
Multiplication of inputs	$1 \rightarrow 100$		
	$100 \times 2 \rightarrow 100$		
Fully connected	$100 \rightarrow 6272$	ReLU	
Reshape	$7 \times 7 \times 128$		
Up-sampling	$14 \times 14 \times 128$		
Convolution	$14 \times 14 \times 128$	ReLU	3x3
Up-sampling	$28 \times 28 \times 128$		
Convolution	$28 \times 28 \times 64$	ReLU	3x3
Convolution	$28 \times 28 \times 1$	tanh	3x3

Discriminator

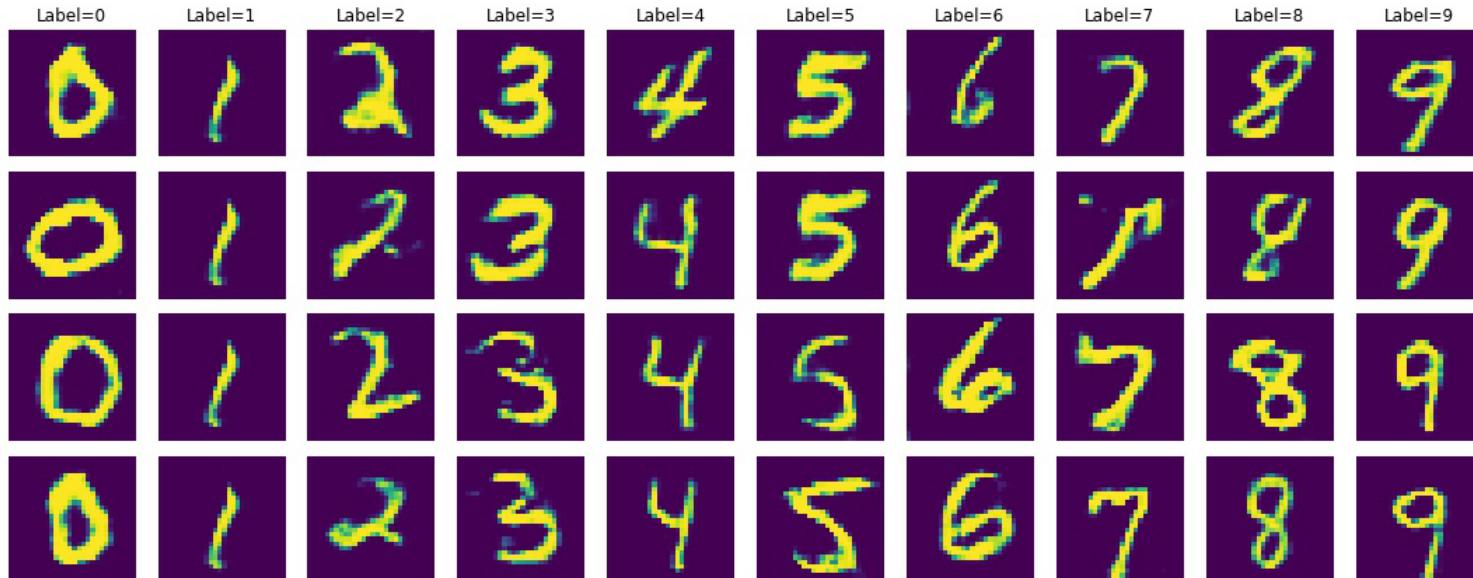
Layer	Size out	Activation	Kernel
Input	$28 \times 28 \times 1$		
Convolution	$28 \times 28 \times 16$	LeakyReLU	3x3
Down-sampling	$14 \times 14 \times 16$		
Dropout			
Convolution	$14 \times 14 \times 32$	LeakyReLU	3x3
Down-sampling	$7 \times 7 \times 32$		
Zero-padding	$8 \times 8 \times 32$		
Dropout			
Convolution	$8 \times 8 \times 64$	LeakyReLU	3x3
Down-sampling	$4 \times 4 \times 64$		
Dropout			
Convolution	$4 \times 4 \times 128$	LeakyReLU	3x3
Dropout			
Output	$2048 \rightarrow 1$	Sigmoid	
Fully connected	$2048 \rightarrow 10$	Softmax	

- Mnist digits as real samples
 - 60000 real samples
- Generator
 - Adam optimizer
 - Learning rate=0.0002
- Generator inputs:
 - Random normal between 0 and 1
 - Label of the digit
- Discriminator
 - Adam optimizer
 - Learning rate=0.0002
- Discriminator outputs:
 - Determining if the input is real or fake
 - Classifying the input image
- 30 epochs

AC-GAN: generating digits

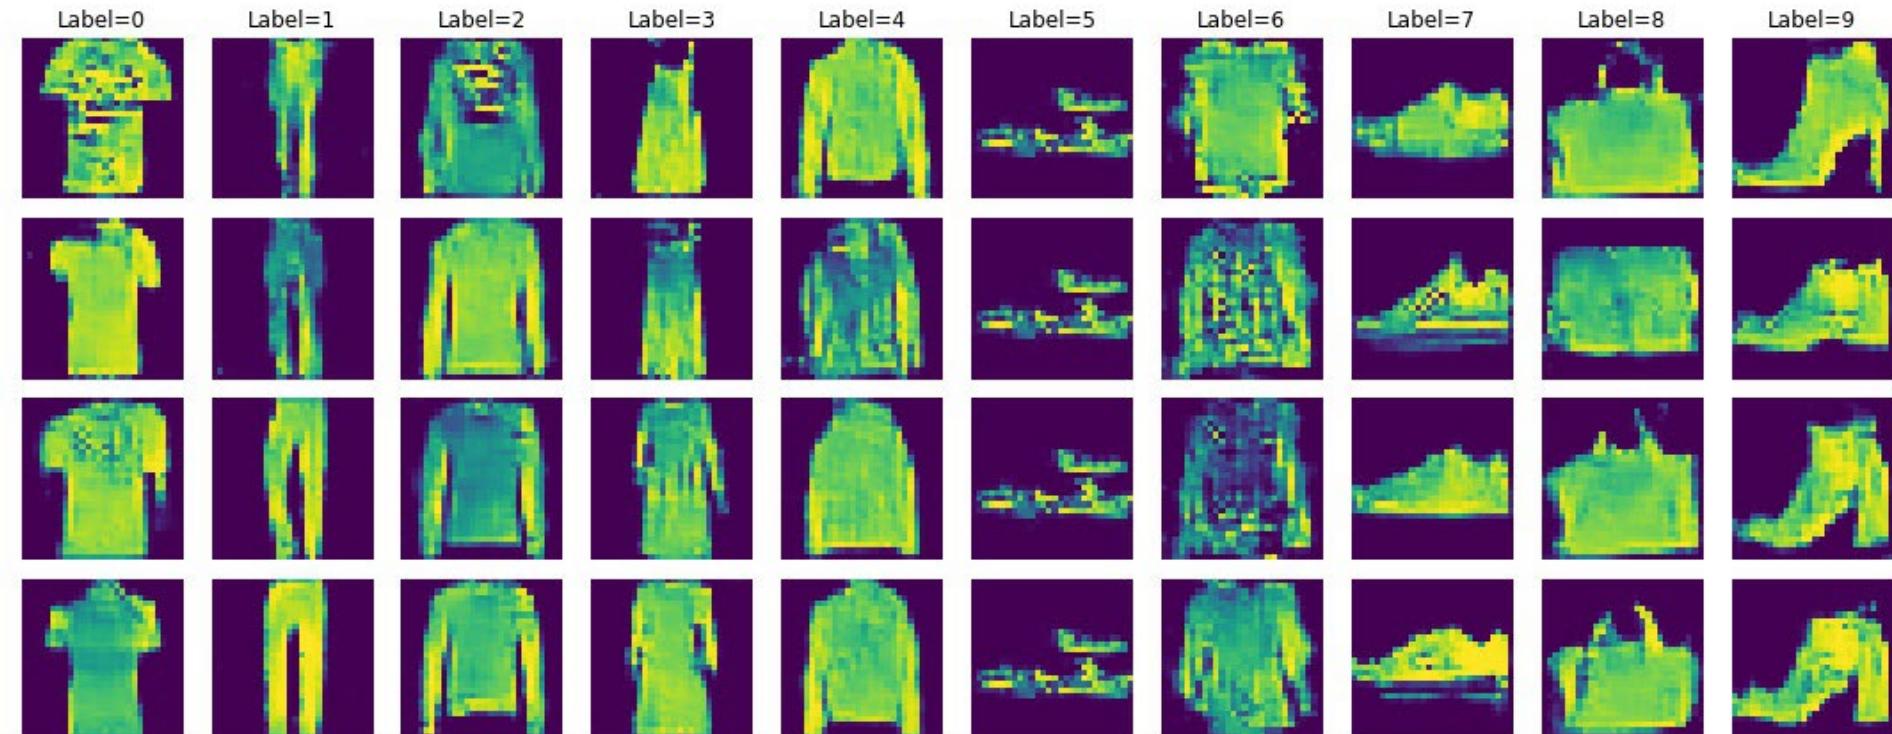


Examples of generated digits

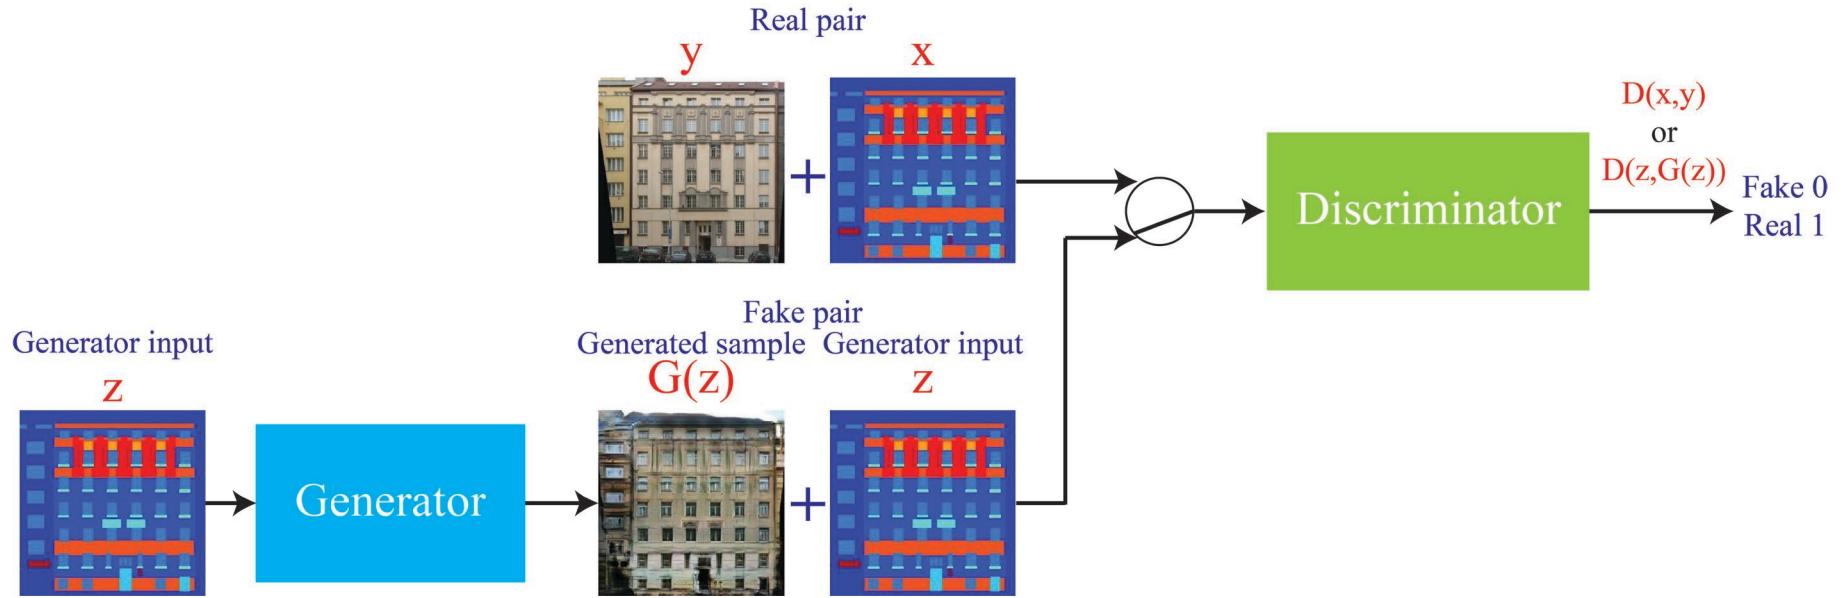


AC-GAN: generating images like mnist-fashion

Examples of generated garments



GAN: Image-to-Image translation



Cost function:

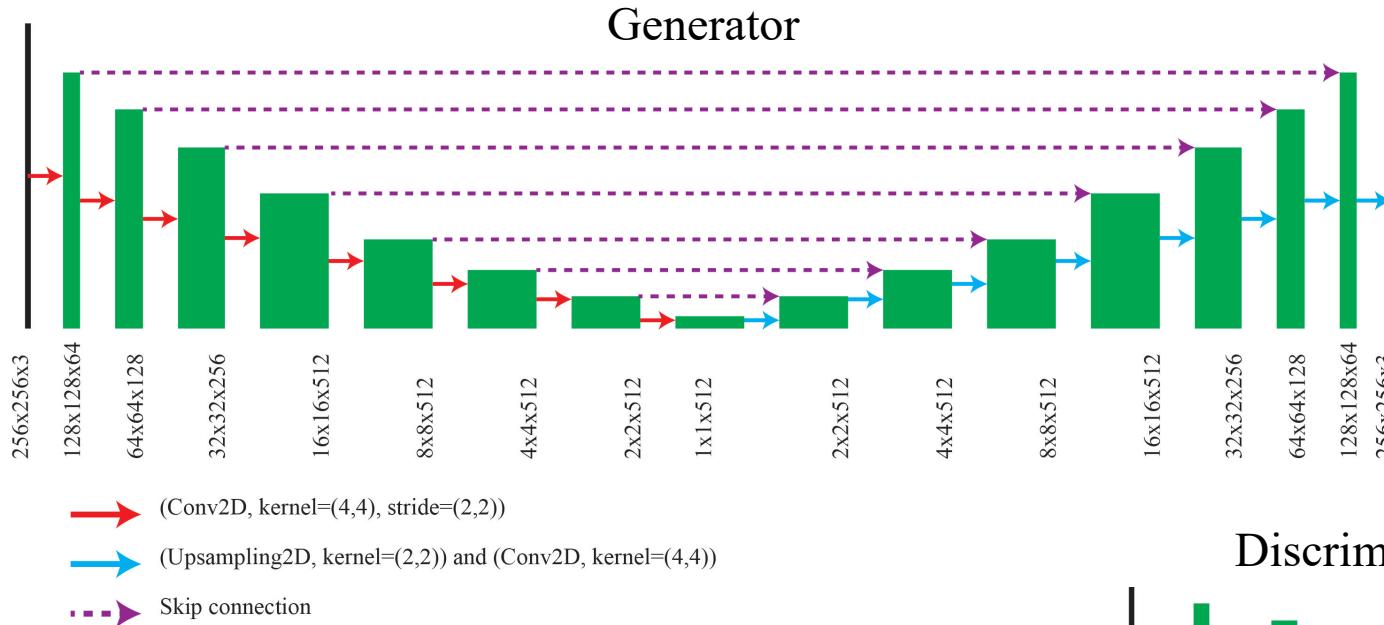
$$C = \sum_{x,y} [\log(D(x,y))] + \sum_z [\log(1 - D(z, G(z)))]$$

$$C_{L1} = \sum_{z,y} \|y - G(z)\|_1$$

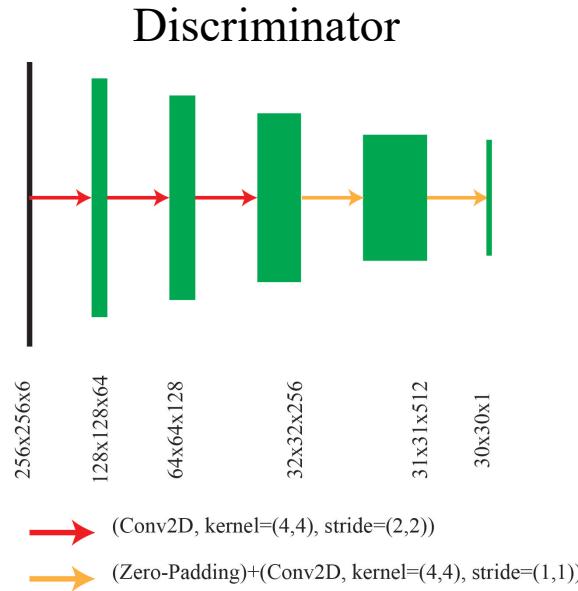
- Discriminator maximizes C
- Generator minimizes $C + C_{L1}$

- x and z are sketches of buildings
- y is the target image
- $G(z)$ is the generated data
- $D(x,y)$ is the discriminator output

GAN: Image-to-Image translation

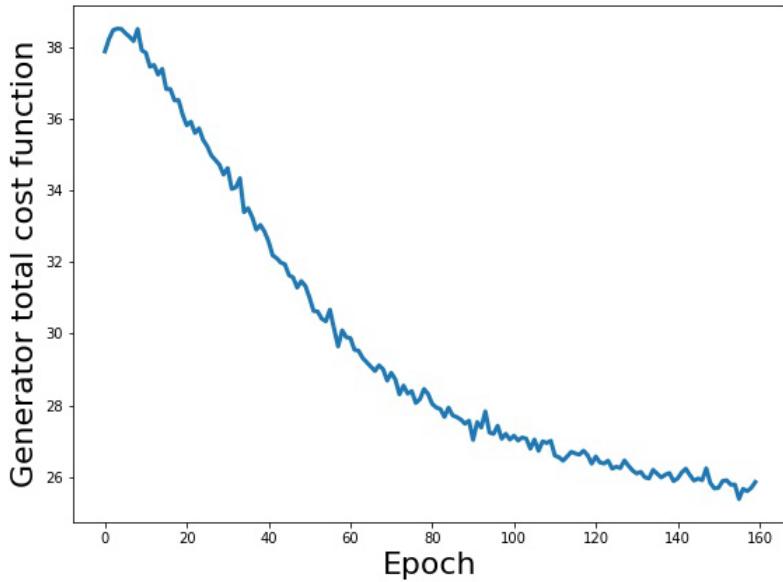


- Generator
 - Adam optimizer
 - Learning rate=0.0002
- Discriminator
 - Adam optimizer
 - Learning rate=0.0002
- 160 epochs

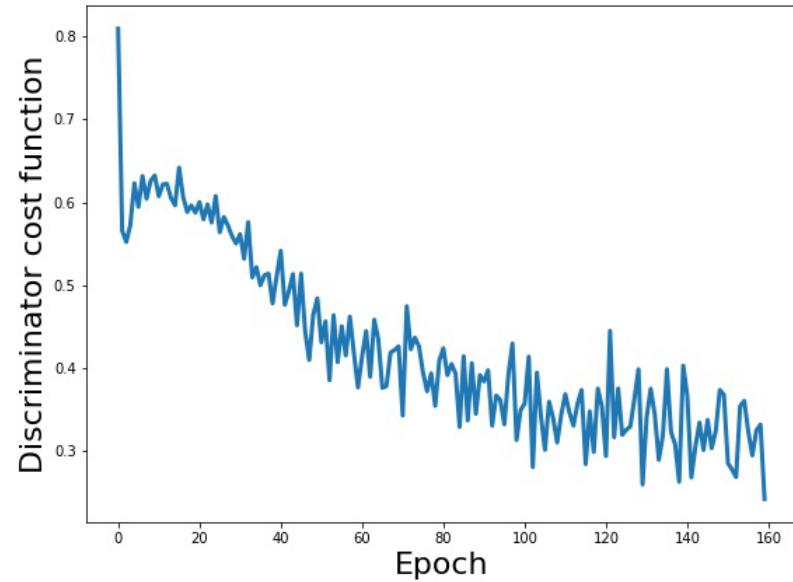


GAN: Image-to-Image translation

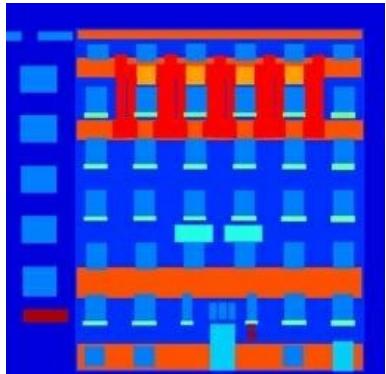
Generator



Discriminator



Generator input



Target

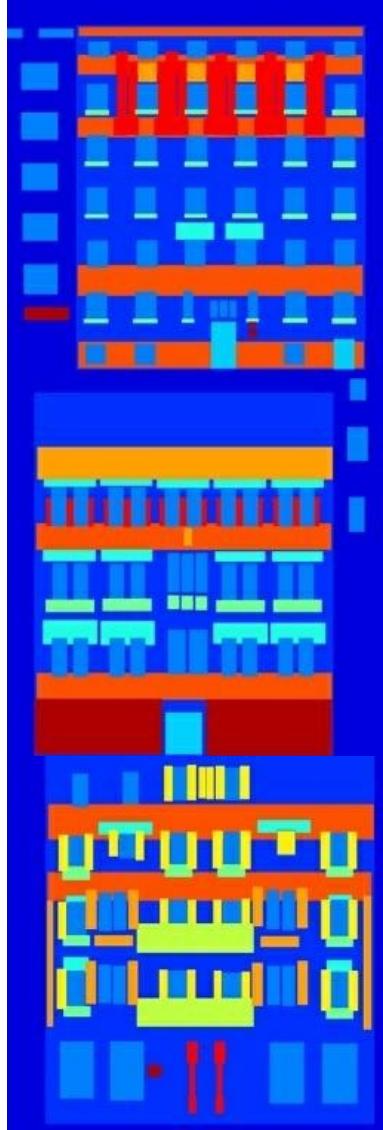


GAN

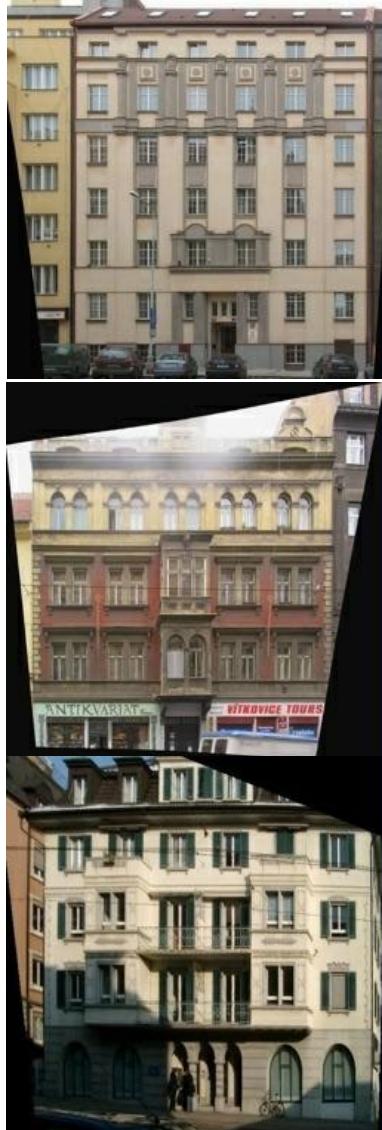


GAN: Image-to-Image translation

Generator input



Target



GAN



U-net

